



Thothica

AN AI STARTUP · DATA ONTOLOGY

AI Systems for Institutions That Run on Records.

We write the ontology, send the agents that fill it, and build the system your staff work in. Then we run it.

Four layers, one company.

5

government portals running our AI layer

115,652

records classified on one of those portals

1,000+

historical works mapped and searchable

10M

words of research made usable


Four Layers, All of Them Ours

This is normally four separate suppliers, plus a systems integrator hired to connect them. We do all four parts ourselves.

L1 ONTOLOGY

We write down what your data means


Our software reads the systems you already run and writes the first draft: what each record is, how it links to the others, and who has to sign off. We set the rules it follows and check what it produces. You end up with one document your staff can read and your software can execute.



L2 AGENTS AND ENGINEERS

We sit with your experts, in software or in person


Much of what decides an outcome was never written down anywhere. It is in the head of the officer who has done the job for eleven years. Forward deployed agents are software that works through the systems your people already use. Forward deployed engineers are our own people, in your building, writing the same structure by hand where software cannot reach. Which of the two we send depends on the engagement, and on some it is both.



L3 APPLICATION

We build the system your staff work in

Either we build it, or we put the layer inside the software you already use. Either way your staff get an answer they can act on, with the source record attached to it.



L4 OPERATION

We run it, and we keep running it


Five government portals, a hundred-year-old archive, a publishing line and a newsroom engine are live on this stack right now, and we operate them. We do not hand over the code and leave.



What an Ontology Is

Ask two departments what a pending case is and you will get two answers. Both are right, because each was defined for a different purpose years ago, and neither definition was ever written down. So the two produce different numbers and nobody can say which one is correct.

An ontology closes it. It is a plain document that says what each thing is, which one is the real number, and who decides.



Those three parts together are the ontology. It is kept in one place, in one structure, and both your staff and your software read it.

WHY THE METHOD REPEATS

Four unrelated industries. No two of them stored their data the same way. We rebuilt the same layer against each one, and the method did not change.

Five Portals, One Read-Only AI Layer

A state government ran five portals on different technology, each with its own idea of how to store the same facts. Officials could see their own screens. Nobody could ask a question across the records.

115,652

records classified on one portal

5

live portals carrying the layer

4

different technologies underneath

0

writes back into any source system

BEFORE

How it worked

A question goes to the officer who knows. They export a spreadsheet, clean it by hand, and produce a figure nobody else can reproduce.

AFTER

How it works now

A question goes to the layer. It answers from that portal's own rows, and every figure opens the records it was built from.

HOW IT RUNS

- 1 **Catalogue**
For each portal, what its data measures and every legitimate way it can be broken down: by district, by scheme, by month, by status.
- 2 **Classify**
Free-text fields sorted into a fixed set so they can be counted. On one portal that was 115,652 separate decisions about what a piece of text meant.
- 3 **Answer**
An official asks in plain language. The answer is assembled from that portal's own rows.
- 4 **Trace**
Every figure opens the records it was built from, so a number can be defended in the room.

THE RULE UNDERNEATH

For each portal we first wrote down what its data measures and every way it can legitimately be broken down, then classified the free-text fields into a fixed set so they could be counted. The layer holds read permission only. It cannot alter a record even if it is asked to.

A Century of Writing, Searchable by Who Said What

A think tank held the written record of a tradition going back to the 1850s: more than two thousand PDFs, most of them photographs of printed pages, on a site that was slow to load and impossible to search.

1,000+

historical works catalogued

506

writers resolved to one identity each

5

languages read across several scripts

2

declared trust tiers

BEFORE

How it worked

Two thousand PDFs of photographed pages. The collection existed; the knowledge in it did not.

AFTER

How it works now

Ask who wrote about whom, on what subject and in what words. Every relation carries the passage that proves it.

HOW IT RUNS

1 Read

Photographs of printed pages turned into text across five languages and several scripts.

2 Resolve

506 writers matched to one identity each, across scripts, spellings and initials.

3 Relate

Who wrote what, who is the subject of what, who is mentioned inside what, who argued with whom.

4 Prove

Every relation carries the verbatim passage, checked as a substring of the source text.

THE RULE UNDERNEATH

Every relation carries the verbatim passage that proves it, checked as a substring of the source text. A relation whose quotation fails that check is discarded rather than kept with a warning, so a fabricated quotation cannot survive storage.

Graphic Stories at Scale, Every Line Traced to a Source

One of India's largest publishers wanted graphic stories faster than any studio could make them, and accurate enough to put a living person's words in a speech bubble.

10M

words in the research corpus

70+

books converted and chaptered

220+

interviews transcribed

5

publishing lines on one pipeline

BEFORE

How it worked

Pace and accuracy pull against each other, and provenance is a promise somebody has to keep by hand.

AFTER

How it works now

Provenance is a mechanical gate. Every beat points at a source line, and the gate resolves each pointer before the script proceeds.

HOW IT RUNS

- 1 **Corpus**
70+ books and 220+ interviews, split chapter per file so a claim can point at an exact line.

- 2 **Script**
Every beat in the finished script carries a pointer to the source line behind it.

- 3 **Gate**
The gate resolves each pointer before the script proceeds. A broken pointer stops the book.

- 4 **Art**
Characters and pages generated against locked visual specifications, shipped press-ready.

THE RULE UNDERNEATH

Sources are split chapter per file so a claim can point at an exact file and line. Art is generated against locked visual specifications per figure, and ships both as press-ready separated colour and as editable pages in the book's own lettering.

A Newsroom That Runs Unattended and Publishes Nothing

A sports publisher needed coverage on a clock that never stops. Hiring for that clock is expensive and still leaves the small hours uncovered.

0

credentials that can publish

1

manifest every consumer reads from

3

kinds of taxonomy dimension

1

human decision left in place

BEFORE

How it worked

Continuous coverage means hiring for a clock that never stops, and the small hours stay uncovered anyway.

AFTER

How it works now

The engine works the clock and stops at an editorial queue. It holds no credential that can publish.

HOW IT RUNS

- 1 **Read**
Live match data, the news wire and fan forums, on a clock that never stops.
- 2 **Plan**
The engine plans its own slate against a taxonomy declared in one manifest.
- 3 **Write**
In the publisher's voice, against a written style specification, checked against sources it fetches.
- 4 **Stop**
Everything lands in a private queue. An editor decides what runs.

THE RULE UNDERNEATH

The engine reads live signals, plans its own slate, writes against a written style specification and corroborates its claims against sources it actually fetches. Then it stops. On its worst day it fills a queue rather than embarrassing the masthead.

Handwritten Registers Turned Into Verified Rows

An academic olympiad receives registrations on paper: a row per student, a tick per subject, a handwritten total at the bottom. Thousands of sheets a season, every one typed up and checked by hand.

2

independent readings of each total

3

verdicts, one of them silent

0

sums done by the model

100%

of conflicts reported, never repaired

BEFORE

How it worked

Every sheet typed up by one person, then checked by a second, because the data becomes a child's exam entry.

AFTER

How it works now

A model reads the page and code does every sum. Anything it cannot verify comes back marked unverified, never marked correct.

HOW IT RUNS

1 Transcribe

A vision model reads what is actually on the paper, and does nothing else.

2 Recompute

Code recomputes each class from the marks. The model does no arithmetic at all.

3 Reconcile

Compared against the handwritten total, read independently and pinned so a later repair cannot move it.

4 Verdict

Verified, conflicted, or unverified. A sheet with nothing to check against is never called correct.

THE RULE UNDERNEATH

The doctrine is that the model does perception and the code does arithmetic. Evidence that confirms a reading must come from a source independent of it, so the same reading is never used to check itself.

Built to Order, by Institution

Grouped by the institution asking. These are capabilities, not finished projects.

ARCHIVES, LIBRARIES, MUSEUMS

Collections That Answer Questions

Manuscripts and oral history captured properly, a public platform over the collection, an open interface any AI can read under your rules, and a picture of who actually cites you.

GOVERNMENT

A Different Use Case per Department

Industry, revenue, procurement, drafting, disaster management, local bodies, constituency offices and public sector oversight. Each has its own problem, and each gets its own first deliverable.

MEDIA AND PUBLISHING

From Research to a Story Ready to Run

Coverage that stops at an editor, a research desk that turns documents into angles, analytics that answer editorial questions, and discourse measured claim by claim.

LEGAL PRACTICES

The Office Keeps Its Judgement

Judgments translated to filing standard under a glossary that is binding, case files structured into positions and authorities, and conflicting authority surfaced across a body of law.

DEFENCE AND SECURITY

Runs Where Nothing May Leave the Building

Doctrine and standing orders made answerable, assessment that keeps source reliability separate from what was reported, and every part able to run fully disconnected.

AI COMPANIES AND LABS

Messy Material Turned Into Training-Grade Structure

Domain corpora typed and given provenance with the licence position recorded, plus evaluation sets a model cannot already have memorised.

A Different Use Case per Department

A revenue department and a disaster management authority share nothing in subject matter and one thing underneath: the rules that decide an outcome are written somewhere no software can read.

DEPARTMENT	WRITTEN DOWN NOWHERE	WHAT THAT COSTS TODAY	FIRST DELIVERABLE
Industry	What each scheme really requires, and in what order	Investors wait on officers who know	An eligibility model behind the single window
Revenue	Which patterns in the register indicate a leak	The same leak is found again every year	Rules as code, each citing its statute
Procurement	What a suspicious award looks like across a body of tenders	Only single files get read	A detection grid over published fields
Drafting	Which instruments contradict which	Conflicts surface in litigation	Obligations typed and compared
Disaster	Who must act on a given warning, and by when	The map lives with experienced officers	Warnings that produce an owned task list
Local bodies	What is pending, and with whom	A diary and a chat thread	Nothing exists without an owner and a date
PSU oversight	Which bodies overlap in mandate	Suspected, never evidenced	A roster where every value opens its source

SEQUENCING

Start with registers the department already owns. Anything that needs another agency's cooperation comes last and is marked optional, because most government data projects stall in a memorandum between two agencies.

The Method, in Public

We publish the method where anyone can inspect it.

OPEN LIBRARY

Falsafa.ai

Our open library of the world's philosophical, classical and religious texts: more than 2,000 works split into more than 1.3 million passages that can each be quoted and cited exactly.

BUILT ON TOP OF IT

The Atlas

A map of the collection drawn from the texts themselves: the people, ideas, places, groups and events across twenty-five centuries, and every time one text cites another, including whether the later author agreed or argued back.

THE RULE UNDERNEATH

It cannot invent a quotation

The software that builds the Atlas is not allowed to write quotations. It can only point at paragraphs, and the words are attached afterwards from the source text.

WHAT IS IN IT TODAY

2,018

works in the library

1.3M

passages addressable

36,472

entities in the Atlas

363,042

verbatim quotations

Thothica works with some of India's largest publishers, with state governments, and with leading think tanks and small businesses.



Thothica

GET IN TOUCH

Tell us your hardest problem.

If the rules that decide an outcome in your organisation are written somewhere no software can read, that is the problem we solve. Send us a description of it and we will tell you what it would take.

THOTHICA PRIVATE LIMITED

hello@thothica.com

thothica.com · New Delhi, India

Clients are described rather than named in this document. Named evidence is available under a confidentiality undertaking.